

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

train_data=pd.read_csv('train.csv')
```

```
train_data
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377
X378 \														
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1
0														
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0
0														
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0
0														
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0
0														
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0
0														
...
...														
4204	8405	107.39	ak	s	as	c	d	aa	d	q	...	1	0	0
0														
4205	8406	108.77	j	o	t	d	d	aa	h	h	...	0	1	0
0														
4206	8412	109.22	ak	v	r	a	d	aa	g	e	...	0	0	1
0														
4207	8415	87.48	al	r	e	f	d	aa	l	u	...	0	0	0
0														
4208	8417	110.85	z	r	ae	c	d	aa	g	w	...	1	0	0
0														
	X379	X380	X382	X383	X384	X385								
0	0	0	0	0	0	0								
1	0	0	0	0	0	0								
2	0	0	1	0	0	0								
3	0	0	0	0	0	0								
4	0	0	0	0	0	0								
...								
4204	0	0	0	0	0	0								
4205	0	0	0	0	0	0								
4206	0	0	0	0	0	0								
4207	0	0	0	0	0	0								
4208	0	0	0	0	0	0								

```
[4209 rows x 378 columns]
```

```
train_data.isna().sum(0)
```

```
ID      0
y        0
X0       0
X1       0
X2       0
..
X380     0
X382     0
X383     0
X384     0
X385     0
Length: 378, dtype: int64
```

```
train_data.describe()
```

	ID	y	X10	X11	X12 \
count	4209.000000	4209.000000	4209.000000	4209.0	4209.000000
mean	4205.960798	100.669318	0.013305	0.0	0.075077
std	2437.608688	12.679381	0.114590	0.0	0.263547
min	0.000000	72.110000	0.000000	0.0	0.000000
25%	2095.000000	90.820000	0.000000	0.0	0.000000
50%	4220.000000	99.150000	0.000000	0.0	0.000000
75%	6314.000000	109.010000	0.000000	0.0	0.000000
max	8417.000000	265.320000	1.000000	0.0	1.000000

	X13	X14	X15	X16	X17
... \					
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
... mean	0.057971	0.428130	0.000475	0.002613	0.007603
... std	0.233716	0.494867	0.021796	0.051061	0.086872
... min	0.000000	0.000000	0.000000	0.000000	0.000000
... 25%	0.000000	0.000000	0.000000	0.000000	0.000000
... 50%	0.000000	0.000000	0.000000	0.000000	0.000000
... 75%	0.000000	1.000000	0.000000	0.000000	0.000000
... max	1.000000	1.000000	1.000000	1.000000	1.000000
...					

	X375	X376	X377	X378	X379
\					
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000

mean	0.318841	0.057258	0.314802	0.020670	0.009503
std	0.466082	0.232363	0.464492	0.142294	0.097033
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	X380	X382	X383	X384	X385
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.007603	0.001663	0.000475	0.001426
std	0.089524	0.086872	0.040752	0.021796	0.037734
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

[8 rows x 370 columns]

If for any column(s), the variance is equal to zero, then you need to remove those variable(s).

```
variance=pow(train_data.drop(columns={'ID','y'}).std(),2).to_dict()
```

```
for key,value in variance.items():
    if(value==0):
        print("Name=",key)
```

```
Name= X11
Name= X93
Name= X107
Name= X233
```

```
Name= X235
Name= X268
Name= X289
Name= X290
Name= X293
Name= X297
Name= X330
Name= X347
```

```
# Now we Will drop this columns
```

```
train_data=train_data.drop(columns={'X11','X93','X107','X233','X235','X268','X289','X290','X293','X297','X330','X347'})
```

```
train_data.shape
```

```
(4209, 366)
```

```
train_data.isnull().sum().any()
```

```
False
```

```
# creating dependent and independent variables
```

```
train_data_feature=train_data.drop(columns={'ID','y'})
```

```
train_data_target=train_data.y
```

```
train_data_feature.shape
```

```
(4209, 364)
```

```
train_data_target.shape
```

```
(4209,)
```

```
train_data_feature.describe(include='object')
```

	X0	X1	X2	X3	X4	X5	X6	X8
count	4209	4209	4209	4209	4209	4209	4209	4209
unique	47	27	44	7	4	29	12	25
top	z	aa	as	c	d	w	g	j
freq	360	833	1659	1942	4205	231	1042	277

```
# So we Got the columns which are obj Hence we Apply Label Encoding in this
```

```
from sklearn.preprocessing import LabelEncoder
lr=LabelEncoder()
```

```
for i in train_data_feature.columns:
    data_type=train_data_feature[i].dtype
    if data_type=='object':
        train_data_feature[i]=lr.fit_transform(train_data_feature[i])
```

Perform dimensionality reduction.

```
from sklearn.decomposition import PCA
pca=PCA(n_components=0.95)
```

```
train_data_feature_trans=pca.fit_transform(train_data_feature)
```

```
train_data_feature_trans.shape
```

```
(4209, 6)
```

```
# Predict your test_df values using XGBoost.
```

```
# Split the dataset into train set & test set
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(train_data_feature_trans,train_data_target,test_size=.3,random_state=42)
```

```
# Check the Shape Of data
```

```
print(X_train.shape)
```

```
print(X_train.shape)
```

```
print(y_train.shape)
```

```
print(y_test.shape)
```

```
(2946, 6)
```

```
(2946, 6)
```

```
(2946,)
```

```
(1263,)
```

```
# importing the XG Boost( Extreme gradient boosting)
```

```
!pip install xgboost
```

```
Defaulting to user installation because normal site-packages is not writeable
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.7/site-packages (1.0.2)
```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.7/site-packages (from xgboost) (1.4.1)
```

```
Requirement already satisfied: numpy in /usr/local/lib/python3.7/site-packages (from xgboost) (1.21.5)
```

```
WARNING: You are using pip version 22.0.3; however, version 22.3.1 is available.
```

```
You should consider upgrading via the '/usr/local/bin/python3 -m pip install --upgrade pip' command.
```

```
import xgboost as xgb
```

```
# Train the model
```

```
xgb_reg=xgb.XGBRegressor()
```

```
model=xgb_reg.fit(X_train,y_train)
```

```
# Prediction
```

```
y_pred=model.predict(X_test)
```

```
# Evaluation
```

```
from sklearn.metrics import mean_squared_error
```

```
print("RMSE IS : ",np.sqrt(mean_squared_error(y_pred,y_test)))
```

```
RMSE IS :  11.813608308644344
```

```
#Saving the model
```

```
import joblib
```

```
joblib.dump(model,'xgbmodel.pkl')
```

```
['xgbmodel.pkl']
```

```
# oad the Model by using loaded model
```

```
loaded_model=joblib.load('xgbmodel.pkl')
```

```
print('Model loaded successfully')
```

```
Model loaded successfully
```

```
# Now On test data
```

```
test_data=pd.read_csv('test.csv')
```

```
test_data=test_data.drop(columns={'X11','X93','X107','X233','X235','X268','X289','X290','X293','X297','X330','X347'})
```

```
# Check For The null Values
```

```
test_data.isnull().sum().any()
```

```
False
```

```
test_data_feature=test_data.drop(columns={'ID'})
```

```
test_data_feature.shape
```

```
(4209, 364)
```

```
# Apply label encoder.
```

```
for i in test_data_feature.columns:
```

```
    data_type=test_data_feature[i].dtype
```

```
    if data_type=='object':
```

```
        test_data_feature[i]=lr.fit_transform(test_data_feature[i])
```

```
test_data_feature
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377
X378	\													
0	21	23	34	5	3	26	0	22	0	0	...	0	0	0
1														
1	42	3	8	0	3	9	6	24	0	0	...	0	0	1
0														
2	21	23	17	5	3	0	9	9	0	0	...	0	0	0

```

1
3      21  13  34   5   3  31  11  13   0   0 ...   0   0   0
1
4      45  20  17   2   3  30   8  12   0   0 ...   1   0   0
0
...   ..  ..  ..   ..  ..  ..  ..  ..  ...  ...  ...  ...  ...
...
4204   6   9  17   5   3   1   9   4   0   0 ...   0   0   0
0
4205  42   1   8   3   3   1   9  24   0   0 ...   0   1   0
0
4206  47  23  17   5   3   1   3  22   0   0 ...   0   0   0
0
4207   7  23  17   0   3   1   2  16   0   0 ...   0   0   1
0
4208  42   1   8   2   3   1   6  17   0   0 ...   1   0   0
0

```

```

      X379 X380 X382 X383 X384 X385
0         0    0    0    0    0    0
1         0    0    0    0    0    0
2         0    0    0    0    0    0
3         0    0    0    0    0    0
4         0    0    0    0    0    0
...      ...  ...  ...  ...  ...  ...
4204      0    0    0    0    0    0
4205      0    0    0    0    0    0
4206      0    0    0    0    0    0
4207      0    0    0    0    0    0
4208      0    0    0    0    0    0

```

[4209 rows x 364 columns]

Perform dimensionality reduction.

```

from sklearn.decomposition import PCA
pca=PCA(n_components=.95)

```

```

test_data_feature_trans=pca.fit_transform(train_data_feature)

```

```

test_data_feature_trans.shape

```

(4209, 6)

Predict your test_df values using XGBoost.

```

test_pred=loaded_model.predict(test_data_feature_trans)

```

```

test_pred

```

```
array([ 92.21506 ,  90.38116 ,  73.274506, ..., 109.36014 ,
        88.40276 ,
        99.617065], dtype=float32)
```